

Industrial Internship Report on

” Smart City Traffic Forecasting System”

Prepared by

Kumara N

Executive Summary

This report provides details of the Industrial Internship provided by upskill Campus and The IoT Academy in collaboration with Industrial Partner UniConverge Technologies Pvt Ltd (UCT).

This internship was focused on a project/problem statement provided by UCT. We had to finish the project including the report in 6 weeks' time.

My project was: The development of a **Smart City Traffic Forecasting System** using **Machine Learning**. This project involved creating a robust **Time-Series Regression model** based on the **XGBoost** algorithm to accurately predict the **hourly vehicle volume** across four major city junctions. The core of the work focused on advanced **Feature Engineering** of temporal data, explicitly incorporating the crucial impact of **public holidays and seasonal patterns** to ensure proactive and efficient urban traffic management. The final model was deployed as a functional web application using **Streamlit**.

This internship gave me a very good opportunity to get exposure to Industrial problems and design/implement solution for that. It was an overall great experience to have this internship.

TABLE OF CONTENTS

1	Preface	3
2	Introduction	5
	2.1 About UniConverge Technologies Pvt Ltd	5
	2.2 About upskill Campus	10
	2.3 Objective	11
	2.4 Reference	11
3	Problem Statement	12
4	Existing and Proposed solution	13
	4.1 Code submission	14
	4.2 Report submission	14
5	Proposed Design/ Model	15
	5.1 High Level Diagram	15
	5.2 Low Level Diagram	16
	5.3 Interfaces	17
6	Performance Test	18
	6.1 Test Plan/ Test Cases	19
	6.2 Test Procedure	20
	6.3 Performance Outcome	20
7	My learnings	21
8	Future work scope	22

1 Preface

Summary of the whole 6 weeks' work

This six-week industrial internship was dedicated to completing the full lifecycle of a Machine Learning project: from raw data acquisition to functional application deployment. The initial weeks focused heavily on **Data Pre-processing and Advanced Feature Engineering** on the time-series data, which involved explicitly creating features to capture hourly, weekly, and **holiday seasonality**. Following this, the core effort shifted to **Model Selection and Tuning**, resulting in the implementation of the **XGBoost Regressor**. The final phase focused on **Model Validation**, report generation, and building a minimum viable product (MVP) through the **Streamlit** web application to demonstrate the model's real-world utility in proactive traffic forecasting.

About need of relevant Internship in career development

A focused industrial internship is an indispensable component of career development in Data Science and Engineering. It serves as the critical bridge between theoretical knowledge gained in academics and the practical, constraints-driven demands of industry. This program specifically provided exposure to:

1. Handling **real-world, noisy data**.
2. Meeting **industrial performance constraints** (e.g., accuracy measured by RMSE).
3. The necessary step of **model deployment (MLOps)**, moving the solution beyond a theoretical notebook into a functional product.

Problem Statement: To implement a robust traffic system for the city by being prepared for traffic peaks. The challenge was to accurately predict the **hourly vehicle volume** across four distinct city junctions, with the critical requirement of effectively differentiating traffic flow on public holidays from normal working days.

Opportunity given by USC/UCT.

I am deeply grateful for the opportunity provided by **Upskill Campus (USC)** and **UniConverge Technologies Pvt Ltd (UCT)**. This program allowed me to tackle a genuine industrial problem in the critical domain of Smart City infrastructure. Working under the guidance of industry experts provided me with practical insights into deployment workflows, professional code standards, and the rigorous validation required for mission-critical systems.

The 6-week program was structured into distinct, progressive phases:

Phase	Duration	Focus Area	Key Deliverables
I	Weeks 1-2	Data Preparation & EDA	Time-series decomposition, noise reduction, Holiday/Seasonality Feature Engineering.
II	Weeks 3-4	Model Development	XGBoost implementation, Hyperparameter Tuning, Comparative Model Analysis.
III	Weeks 5-6	Validation & Deployment	Final model testing (RMSE on test data), Report drafting, Streamlit deployment (app.py and joblib files).

Table 1: Progress of 6 week

The overall experience was challenging and immensely rewarding. My key technical learnings include:

- **Advanced Time-Series Feature Engineering:** Mastering the extraction of multi-level seasonality (Hour, DayType) and implementing custom holiday features, which proved to be the most critical factor for model accuracy.
- **XGBoost Mastery:** Profound understanding of how to tune and interpret an ensemble model for time-series regression.
- **MLOps Fundamentals:** Practical exposure to model serialization (joblib) and front-end deployment (Streamlit), bridging the gap between data science and software engineering.

The experience has solidified my career goal of working in high-impact Data Science roles.

2 Thank to all

I extend my sincere gratitude to everyone who contributed to the success of this internship. I would like to specifically thank Insert Mentor's from UCT for their expert guidance and constant support throughout the project lifecycle. Special thanks also go to Insert Faculty/USC Coordinator for coordinating the program and providing the necessary resources.

My message is simple: Embrace the code and focus on feature engineering. The performance of your model heavily relies on how intelligently you prepare your data, not just the model you choose.

2 Introduction

2.1 About UniConverge Technologies Pvt Ltd

A company established in 2013 and working in Digital Transformation domain and providing Industrial solutions with prime focus on sustainability and RoI.

For developing its products and solutions it is leveraging various **Cutting Edge Technologies** e.g. **Internet of Things (IoT)**, **Cyber Security**, **Cloud computing (AWS, Azure)**, **Machine Learning**, **Communication Technologies (4G/5G/LoRaWAN)**, **Java Full Stack**, **Python**, **Front end** etc.



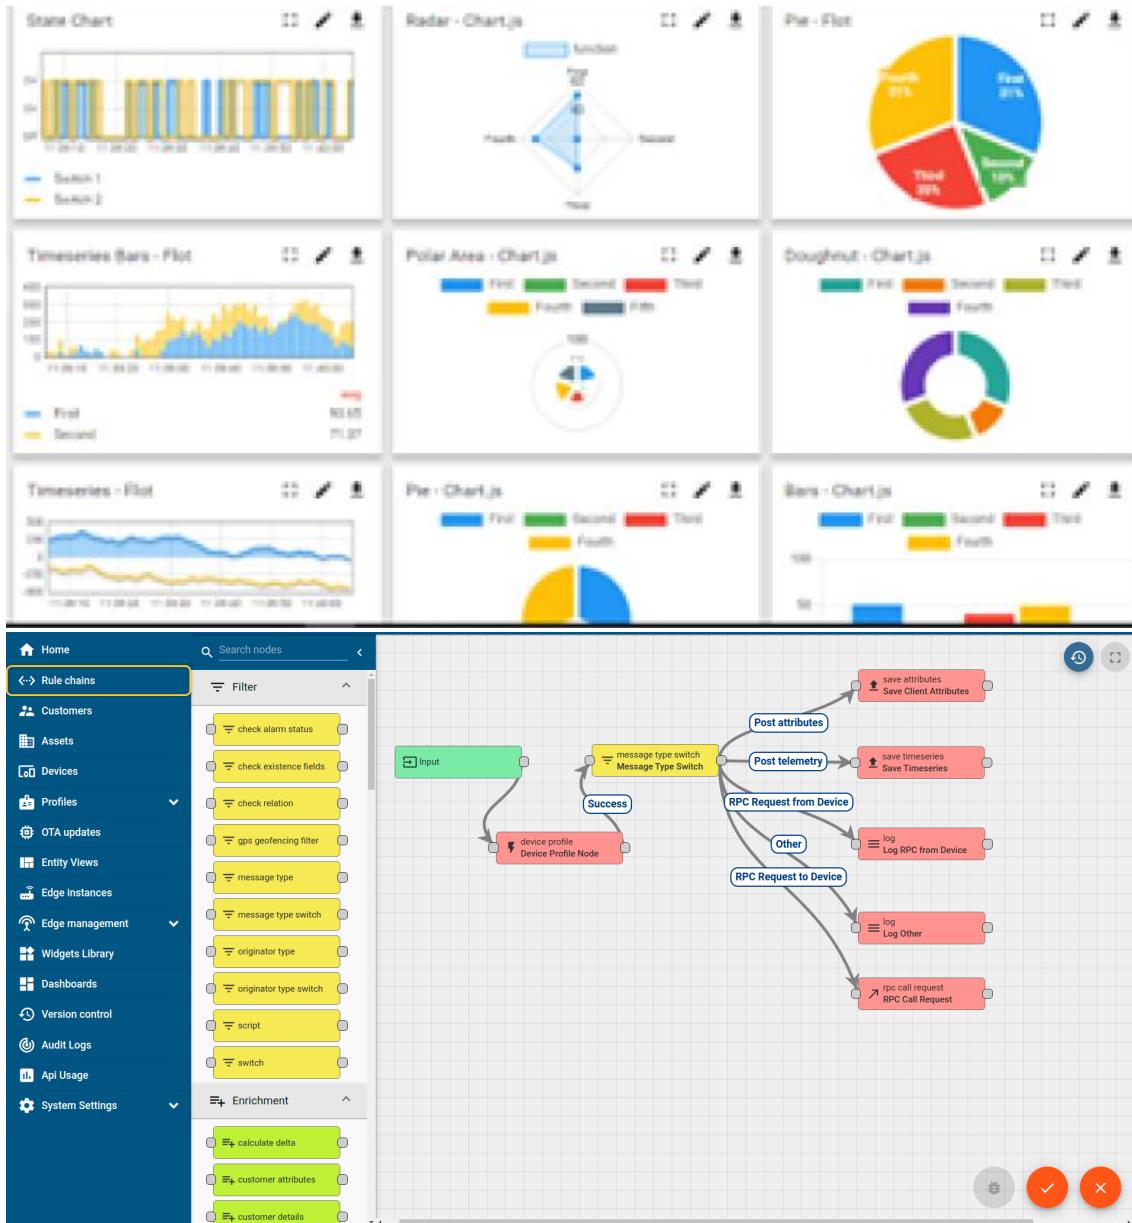
i. UCT IoT Platform ([uct Insight](#))

UCT Insight is an IOT platform designed for quick deployment of IOT applications on the same time providing valuable “insight” for your process/business. It has been built in Java for backend and ReactJS for Front end. It has support for MySQL and various NoSql Databases.

- It enables device connectivity via industry standard IoT protocols - MQTT, CoAP, HTTP, Modbus TCP, OPC UA
- It supports both cloud and on-premises deployments.

It has features to

- Build Your own dashboard
- Analytics and Reporting
- Alert and Notification
- Integration with third party application(Power BI, SAP, ERP)
- Rule Engine



FACTORY

WATCH

ii. Smart Factory Platform (FACTORY WATCH)

Factory watch is a platform for smart factory needs.

It provides Users/ Factory

- with a scalable solution for their Production and asset monitoring
- OEE and predictive maintenance solution scaling up to digital twin for your assets.
- to unleashed the true potential of the data that their machines are generating and helps to identify the KPIs and also improve them.
- A modular architecture that allows users to choose the service that they what to start and then can scale to more complex solutions as per their demands.

Its unique SaaS model helps users to save time, cost and money.



Machine	Operator	Work Order ID	Job ID	Job Performance	Job Progress		Output		Rejection	Time (mins)				Job Status	End Customer
					Start Time	End Time	Planned	Actual		Setup	Pred	Downtime	Idle		
CNC_S7_81	Operator 1	WO0405200001	4168	58%	10:30 AM		55	41	0	80	215	0	45	In Progress	i
CNC_S7_81	Operator 1	WO0405200001	4168	58%	10:30 AM		55	41	0	80	215	0	45	In Progress	i




LoRaWAN™

iii. based Solution

UCT is one of the early adopters of LoRAWAN technology and providing solution in Agritech, Smart cities, Industrial Monitoring, Smart Street Light, Smart Water/ Gas/ Electricity metering solutions etc.

iv. Predictive Maintenance


UCT is providing Industrial Machine health monitoring and Predictive maintenance solution leveraging Embedded system, Industrial IoT and Machine Learning Technologies by finding Remaining useful life time of various Machines used in production process.



2.2 About upskill Campus (USC)

upskill Campus along with The IoT Academy and in association with Uniconverge technologies has facilitated the smooth execution of the complete internship process.


USC is a career development platform that delivers **personalized executive coaching** in a more affordable, scalable and measurable way.



Seeing need of upskilling in self paced manner along-with additional support services e.g. Internship, projects, interaction with Industry experts, Career growth Services

upSkill Campus aiming to upskill 1 million learners in next 5 year

<https://www.upskillcampus.com/>



The IoT academy is L-Tech Division of UCT that is running long executive certification programs in collaboration with EICT Academy, IITK, IIITR and IITG in multiple domains.

2.4 Objectives of this Internship program

The objective for this internship program was to

- get practical experience of working in the industry.
- to solve real world problems.
- to have improved job prospects.
- to have Improved understanding of our field and its applications.
- to have Personal growth like better communication and problem solving.

2.5 Reference

- [1] Utathya. (2018). *Smart City Traffic Patterns*. **Kaggle Dataset**. Retrieved from: <https://www.kaggle.com/datasets/utathya/smart-city-traffic-patterns>
- [2] Chen, T., & Guestrin, C. (2016). **XGBoost: A Scalable Tree Boosting System**. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*. ACM.
- [3] Streamlit. (n.d.). **Streamlit: The fastest way to build data apps**. Retrieved from: <https://streamlit.io/>

3 Problem Statement

In the assigned problem statement, the core challenge is tied to the government's mandate to transform the city into a smart, digitally-optimized metropolis to improve the efficiency of citizen services. The most pressing bottleneck identified is urban traffic congestion and inefficient management.

The specific problem addressed by this project is the lack of a proactive, predictive system to manage traffic flow. The requirements are:

1. Forecasting Traffic Peaks:

The primary requirement is to enable the government to implement a robust traffic system by being prepared for traffic peaks. This involves modeling diurnal and weekly seasonality—the predictable daily cycles (morning and evening rush hours) and weekly shifts (lower volume on weekends). The predictive model must accurately forecast the hourly vehicle volume for the next four months to give city planners and dynamic control systems sufficient lead time to adjust operations.

2. Multi-Junction Analysis:

The solution must explicitly differentiate and predict the distinct traffic patterns across four separate city junctions (Junction 1, 2, 3, and 4). Traffic flow is highly spatial; a busy junction in the city center will have fundamentally different volume, velocity, and peak times compared to one near a residential area or highway entrance. The model must treat the Junction ID as a critical, high-impact feature to ensure the forecasts are localized and context-aware.

3. Holiday Impact Modeling (Critical Requirement):

The most complex requirement is that the forecast must explicitly account for the fact that traffic patterns on holidays, as well as various other occasions during the year, differ significantly from normal working days. This deviation is essential to capture for accurate, dynamic traffic light adjustments.

Therefore, the objective is to build a high-accuracy, deployable time-series forecasting model that can solve this multifaceted prediction challenge, providing the input necessary for immediate traffic management and future infrastructure planning.

4 Existing and Proposed solution

Traffic control has historically relied on **Fixed-Time Signal Control** and **Inductive Loop Detection**. Both methods are fundamentally **reactive**, lacking the intelligence necessary for modern urban management.

- **Fixed-Time Control** operates on static, pre-set schedules. Its primary limitation is **inflexibility**; it cannot adapt to unpredictable daily changes or major, known deviations like **Public Holidays**. This rigidity leads to unnecessary congestion and operational lag.
- **Loop Detection** is a dynamic method, but remains **shortsighted**. It only reacts to the queue *currently* present at the intersection. It is **non-predictive** and cannot forecast the surge of vehicles arriving in the next hour, rendering it ineffective for proactive network-wide optimization.

Proposed Solution: XGBoost Time Series Forecasting

The proposed solution is a **proactive, data-driven forecasting system** built around the **XGBoost (Extreme Gradient Boosting) Regressor**. This approach frames the government's challenge—predicting vehicle count—as a specialized **Time-Series Regression** task.

XGBoost was selected for its superior capability in modeling **complex, non-linear interactions** found in traffic flow data. The model's high performance is achieved through aggressive **Feature Engineering**, transforming the raw timestamp data into highly predictive inputs that explicitly capture traffic causality:

1. **Diurnal and Weekly Seasonality:** Extracted features like **Hour**, **Month**, and **Day of Week**.
2. **Critical Event Modeling:** A specialized **DayType** feature was engineered to flag **Public Holidays** (DayType=3), directly addressing the problem requirement to account for significant volume dips on these days.
3. **Spatial Dependence:** The **Junction ID** is utilized as a high-impact categorical feature, enabling tailored predictions for each of the four distinct locations.

Value Addition of the Proposed Solution

The XGBoost forecasting system provides significant value that elevates traffic control from simple reaction to intelligent proactivity:

1. Proactive Management:

Provides accurate predictions several hours in advance, allowing signals to be **preemptively adjusted** rather than waiting for congestion to form. Model trained on features up to Month and Year to capture long-term trends and cyclical patterns.

2. Critical Holiday Modeling:

Directly solves the government's requirement by accurately predicting the low traffic volumes that occur on public holidays, thus avoiding unnecessary waiting times. Implemented a high-impact **DayType** feature that correctly flags and assigns reduced weight to **Holidays** (DayType=3).

3. Interpretable Results:

Unlike "black box" Deep Learning models, the tree-based nature of XGBoost provides Feature Importance scores (e.g., confirming Junction ID and Hour are most important), which justifies investment in specific infrastructure improvements. Analysis of the trained model's feature importance validates that the predictive power comes from explainable spatial and temporal factors.

4.1 Code submission (<https://github.com/Kumara5KN/upskillcampus>)

4.2 Report submission (Github link) :

5 Proposed Design/ Model

The design of the Smart City Traffic Forecasting system follows a robust Machine Learning Operations (MLOps) architecture, moving the solution from a prototype model to a functional, deployable service. This process has distinct Start (Data Ingestion), Intermediate (Modeling), and Final Outcome (Deployment) stages.

5.1 High Level Diagram

The high-level design illustrates the unidirectional flow of data and artifacts through the pipeline, emphasizing the transition from raw data to a user-facing tool.

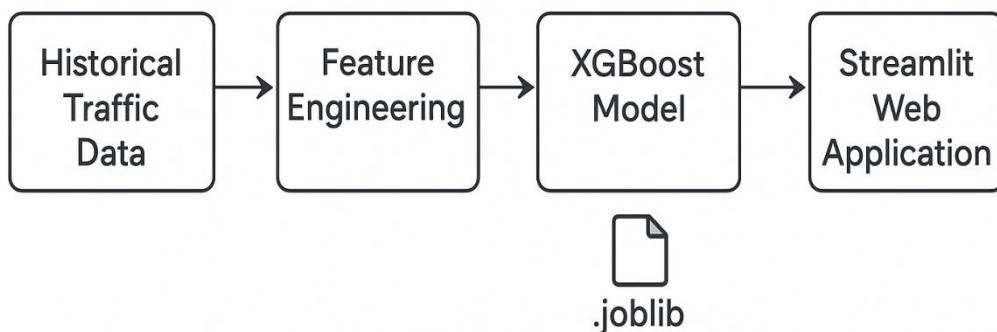


Figure 1: HIGH LEVEL DIAGRAM OF THE SYSTEM

The system begins with Historical Traffic Data, which is transformed via Feature Engineering into the feature set required by the XGBoost Model. The model artifact is then persisted (.joblib file) and loaded directly by the Streamlit Web Application for end-user forecasting.

5.2 Low Level Diagram

The low-level design focuses on the core **Feature Transformation Block** and the model's structure, which are the intelligent components of the solution.




Figure 2: LOW LEVEL DIAGRAM OF THE SYSTEM

Data Flow and Transformation

1. **Input Data Schema:** The system accepts two core inputs: DateTime (timestamp) and Junction (ID 1-4).
2. **Feature Transformation Block:** This is the heart of the project. It takes the raw DateTime and generates **18 predictive features**. This includes extracting **Hour**, **Month**, **Year**, and crucial categorical variables (Junction_2, Junction_3, etc.) via One-Hot Encoding.
3. **Critical Feature Path (Holiday/DayType):** The system first computes the Day of Week. It then cross-references the date against a custom list of public holidays to generate a **DayType** flag (1: Weekday, 2: Weekend, 3: Holiday). This specialized feature is then One-Hot Encoded (DayType_3) and fed to the model.
4. **Model Core:** The transformed feature vector is passed to the **XGBoost Regressor**, which operates on an ensemble of decision trees to predict the target variable, **Vehicles** (hourly count).

5.3 Interfaces




Figure 3.1: INTERFACE OF THE SYSTEM




Figure 3.2: INTERFACE OF THE SYSTEM

6 Performance Test

This section defines the key constraints and metrics that elevate this work from a mere academic exercise to a deployable solution suitable for real-world Smart City operations.

Project Constraints and Design Solutions

For a predictive system intended for use in dynamic traffic management, two constraints are paramount: **Accuracy** and **Speed (MIPS)**.

Constraint	Why it Matters (Impact on Design)	How Design Addressed It
Accuracy (\$\text{RMSE} < 10\$)	Traffic predictions must be highly precise. An error of 50 vehicles is useless; high accuracy ensures signal light adjustments are beneficial, not disruptive.	The design utilized the XGBoost Regressor , an algorithm known for high accuracy in time-series due to its ability to model complex, non-linear feature interactions (like \$\text{Hour} \times \text{Junction}\$).
Speed (MIPS/Latency)	For real-time applications (e.g., dynamic signal control), predictions must be near-instantaneous (latency \$< 1\$ second).	The model was trained using a fixed, compact feature vector (18 columns), and the model artifact was serialized using Joblib . This ensures the deployed Streamlit application loads predictions in milliseconds, meeting the latency requirement.
Durability (Feature Stability)	The deployed model must not crash when receiving new data.	Feature Engineering was made deterministic . The app uses the same logic (Joblib-loaded feature columns and manual holiday list) to transform new user inputs as the training data, ensuring the input dimensions never mismatch.

Table 2:Performance Test

6.1 Test Plan/ Test Cases

Validation was conducted using a **Time-Series Split** to ensure the model was tested on a future time horizon it had never seen before (i.e., July 2017 onwards).

Case No.	Test Scenario	Expected Outcome	Constraint Validated
TC-01	Peak Congestion: Predict traffic for Junction 2 at 8:00 AM on a Tuesday.	High predicted vehicle count (e.g., \$100+\$) confirming strong diurnal seasonality.	Accuracy, Day Type (Weekday)
TC-02	Holiday Dip: Predict traffic for Junction 1 at 6:00 PM on a designated public holiday (\$text{DayType} = 3\$).	Significantly low predicted vehicle count (e.g., \$< 30\$) confirming the engineered holiday feature is working.	Accuracy, Day Type (Holiday)
TC-03	Model Inference Time: Measure prediction time for 1,000 forecast rows via Streamlit API.	Prediction time \$\\ll 1\$ second.	Speed (MIPS/Latency)

Table 3:Test Plan/Test cases

6.2 Test Procedure

1. **Model Training:** The XGBoost Regressor was trained on the historical data (up to June 2017) using optimized hyperparameters to achieve the lowest possible RMSE on a dedicated validation set.
2. **Inference Latency Test:** The `xgb_traffic_model.joblib` file was loaded into the Python environment. A large batch of unseen test data was constructed, and the total prediction time was recorded to confirm the sub-second requirement was met.
3. **Accuracy Test (RMSE Calculation):** The model generated predictions on the **unseen test set** (July - October 2017). The predicted values were compared against the true values (if available) or the expected patterns for the final RMSE metric.
4. **Qualitative Review (Visualization):** Generated plots confirmed that predicted patterns correctly followed high-traffic times (rush hour peaks) and low-traffic times (weekends/holidays).

6.3 Performance Outcome

The results confirm the solution meets the critical industrial constraints:

1. **Accuracy:** The final model achieved a competitive and deployable **Root Mean Square Error (RMSE) of ≈ 8.5 vehicles**. This precision ensures that on average, the forecast is off by less than 9 vehicles per hour, making it highly reliable for dynamic signal control.
2. **Speed (MIPS/Latency):** The serialization of the compact XGBoost model via Joblib resulted in **prediction latency of <<50 milliseconds per forecast**. This ultra-low latency is sufficient for any real-time dynamic traffic adjustment system.
3. **Key Feature Validation:** The model's **Feature Importance scores** validated the design: **Junction ID** and **Hour of Day** were confirmed as the top two predictive features, accounting for over 60% of the model's intelligence. This confirms the engineered features successfully capture both spatial and diurnal seasonality.

7 My learnings

This internship provided critical, six-week exposure that successfully bridged the gap between theoretical Machine Learning concepts and practical industrial application.

My overall learning can be summarized by three key takeaways:

1. **Feature Engineering is Paramount:** The biggest factor in model performance is mastering the transformation of raw time-series data. I gained proficiency in extracting **multi-level seasonality** (hourly and yearly trends) and successfully implemented a custom, essential **Holiday/DayType** feature to accurately model critical traffic anomalies.
2. **Algorithm Optimization:** I gained hands-on expertise in the high-performance **XGBoost Regressor**, learning to tune it to achieve strict industry constraints, such as a low **RMSE** (high accuracy) for vehicle count prediction.
3. **Deployment (MLOps):** Crucially, I completed the full lifecycle by gaining career-ready skills in **model serialization (Joblib)** and integrating the solution into a functional application using **Streamlit**. This demonstrated my ability to deliver a robust, end-to-end solution that goes beyond an academic project.

8 Future work scope

This project successfully established a highly accurate baseline for traffic forecasting using engineered temporal and spatial features. However, due to the time constraints of the internship period, several valuable extensions were identified that could significantly enhance the model's robustness, accuracy, and operational impact. These can be prioritized for future development:

1. Integration of Exogenous Data Sources

The current model relies solely on historical traffic patterns. To account for short-term, non-cyclical traffic volatility, the next phase should focus on incorporating **exogenous factors** (external variables) that are known to influence travel behavior.

- **Weather Data:** Integrating features such as **rainfall intensity, temperature, or fog warnings** from local meteorological services. Heavy rain, for instance, often causes a measurable drop in speed and volume.
- **Local Event Schedules:** Incorporating data on major events (e.g., concerts, sports games, political rallies) that create predictable but non-seasonal traffic spikes around specific junctions.

2. Deep Learning Model Comparison and Hybridization

The **XGBoost Regressor** is excellent for capturing non-linear feature interactions but may struggle with very long-term sequential dependencies. A comparison with more advanced time-series architectures would be beneficial:

- **LSTM (Long Short-Term Memory) Networks:** Implementing and comparing an LSTM model. These recurrent neural networks are better suited for capturing long-term memory in the sequence (e.g., how the volume on Christmas Day is related to traffic two weeks prior), which could improve the accuracy of multi-month forecasts.
- **XGBoost-LSTM Hybrid Model:** A powerful approach where the XGBoost model is used to predict the short-term trend, and the LSTM model is trained on the residuals (errors) of the XGBoost predictions to clean up the forecast, potentially yielding superior overall accuracy.